



# YANFF — YET ANOTHER NETWORK FUNCTION FRAMEWORK LABS

Intel Golang Team

Gregory Shimansky

# Legal Information

Intel technologies' features and benefits depend on system configuration and may require enabled hardware, software or service activation. Learn more at [intel.com](#), or from the OEM or retailer.

No computer system can be absolutely secure.

Software and workloads used in performance tests may have been optimized for performance only on Intel microprocessors.

Performance tests, such as SYSmark and MobileMark, are measured using specific computer systems, components, software, operations and functions. Any change to any of those factors may cause the results to vary. You should consult other information and performance tests to assist you in fully evaluating your contemplated purchases, including the performance of that product when combined with other products. For more complete information visit <http://www.intel.com/performance>.

Tests document performance of components on a particular test, in specific systems. Differences in hardware, software, or configuration will affect actual performance. Consult other sources of information to evaluate performance as you consider your purchase. For more complete information about performance and benchmark results, visit <http://www.intel.com/performance>.

Cost reduction scenarios described are intended as examples of how a given Intel- based product, in the specified circumstances and configurations, may affect future costs and provide cost savings. Circumstances will vary. Intel does not guarantee any costs or cost reduction.

Results have been estimated or simulated using internal Intel analysis or architecture simulation or modeling, and provided to you for informational purposes. Any differences in your system hardware, software or configuration may affect your actual performance.

Intel does not control or audit third-party benchmark data or the web sites referenced in this document. You should visit the referenced web site and confirm whether referenced data are accurate.

Intel, the Intel logo and others are trademarks of Intel Corporation in the U.S. and/or other countries.

\*Other names and brands may be claimed as the property of others.

© 2017 Intel Corporation.

Copyright © 2017, Intel Corporation. All rights reserved.

\*Other names and brands may be claimed as the property of others.

Optimization  
Notice 



# YANFF - Yet Another Network Function Framework

Framework for building performant native network functions

- Open-source project
- Higher level abstractions than DPDK
- Go language: productivity, performance, concurrency, safety
- Network functions are application programs and not virtual machines

## Benefits:

- Easily leverage IA HW capabilities: multi-cores, AES-NI, CAT, QAT, DPDK
- 10x reduction lines of code
- No need to be expert network system programmer
- Similar performance with C
- Take advantage of cloud native deployment: continuous delivery, micro-services, containers

<https://github.com/intel-go/yanff>

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice



3

# Technical Motivation

- Developers need framework to shorten development cycle of VNFs
  - Currently VNFs are monolithic - “virtual appliances” instead of network functions
  - Significant part of VNF is about plumbing. Plumbing VNFs to CommSPs network is an art. Should be abstracted from VNFs
- Lack of stable and unified APIs for VNF control and data plane
- Challenges with access to HW Accelerators in cloud environment.
- Cloud-friendly APIs and designs needed.

## Accelerating transition to from rule-based networking to *imperative networking*

Copyright © 2017, Intel Corporation. All rights reserved.

\*Other names and brands may be claimed as the property of others.

<https://github.com/intel-go/yanff>

Optimization  
Notice

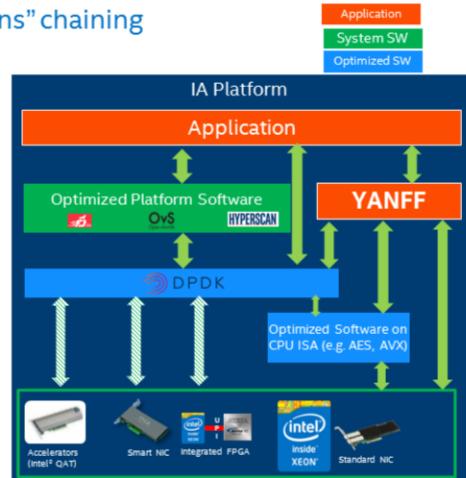


4

VMs have relatively high overhead and memory footprint

# YANFF: Yet Another Network Function Framework

- Simple but powerful abstractions:
  - Flow, Packet
- User builds packet processing graph using “flow functions” chaining
  - SetReceiver -> SetHandler -> SetSender
  - Several predefined possibilities of adding user processing inside packet processing graph
    - Split, Separate, Generate, Handle
- Can leverage predefined functions which parse packets, check ACL rules, etc.
- Run to completion – NFs can be expressed in the flow functions and natural chaining
- Auto-scaling, ease of development
- Zero-copy between NFs
- Flexible incoming flow handling – sources can be anything: network port, memory buffer, remote procedure call, etc.



Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

<https://github.com/intel-go/yanff>



4

## L3 Simple Forwarding Example

```
var L3Rules *rules.L3Rules

func main() {
    flow.SystemInit(16)
    L3Rules = rules.GetL3RulesFromORIG("Forwarding.conf")
    inputFlow := flow.SetReceiver(0)
    outputFlows := flow.SetSplitter(inputFlow, L3splitter, uint(3))
    flow.SetStopper(outputFlows[0])
    for i := 1; i < 3; i++ {
        flow.SetSender(outputFlows[i], uint8(i-1))
    }
    flow.SystemStart()
}

// User defined function for splitting packets
func L3Splitter(currentPacket *packet.Packet) uint {
    currentPacket.ParseL4()
    return rules.L3_ACL_port(currentPacket, L3Rules)
}
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice



# Configuration file for Forwarding

# Source address, Destination address, L4 protocol ID, Source port, Destination port, Output port

111.2.0.0/31	ANY	tcp	ANY	ANY	1
111.2.0.2/32	ANY	tcp	ANY	ANY	Reject
ANY	ANY	udp	3078:3964	56:61020	2

# Exactly The Same Example in DPDK/C

23 SLOC in YANFF vs 2079 in DPDK/C!

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.



# YANFF – Main Architectural Concepts

## Flow

Abstraction without public fields, which is used for pointing connections between **Flow functions**.  
Opened by **Receive / Split / Separate / Counter / Generate**.  
Closed by **Send / Merge / Stop**.

## Packet

High-level representation of network packet. Private field is \*mbuf, public fields are mac / ip / data /etc: pointers to mbuf with offsets (zero copy).  
Is extracted before any **User defined function**. Can be filled after user request by **Packet functions**. Can be checked by **Rule functions**.

## Port

Network door, used in **Receive, Send**.

## Rule

Set of checking rules, used in **User defined functions**.

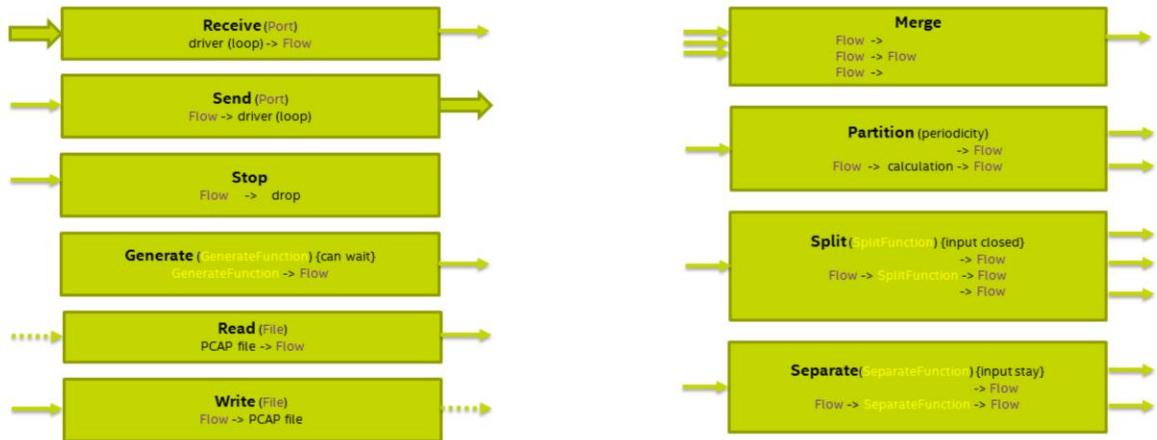
### Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

\*Other names and brands may be claimed as the property of others.



# Building Processing Graph



## Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.



10

# Packet modification functions



## Packet functions

**Parsing packet fields**  
Parse L2 or/and L3 or/and L4 levels

**Initializing packet fields**  
Initialize L2 or/and L3 or/and L4 levels

**Encapsulate / Decapsulate**

## Rule functions

**Create rule**  
Create checking rule from json / config

**Checking packet fields by rule**  
Check L2 or/and L3 or/and L4 levels

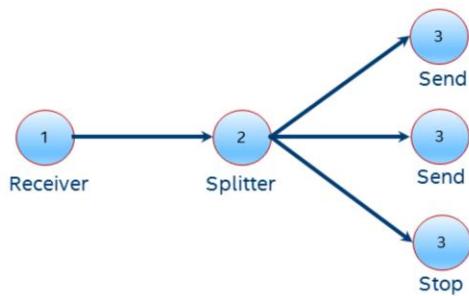
### Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.



11

## Flow Graph Example - Forwarding



Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

\*Other names and brands may be claimed as the property of others.



12

# Let's build some functions!

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.

\*Other names and brands may be claimed as the property of others.



13

## Create test VMs

1.Create and provision two test VMs:

```
$ cd $GOPATH/src/github.com/intel-go/yanff/vagrant  
$ vagrant up
```

2.Open two terminal windows

3.cd to vagrant directory below

4.run “vagrant ssh yanff-VM\_number” to connect to pktgen VM and target VM, e.g.

```
$ vagrant ssh yanff-1      # YANFF test program host  
yanff-1$ bindports          # if ports not bound yet
```

```
$ vagrant ssh yanff-0      # pktgen host  
yanff-0$ bindports          # if ports not bound yet
```

Optimization Notice

Copyright © 2016, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.



14

## Let's try (1 of 11)

Flow graph:

```
package main
import "github.com/intel-go/yanff/flow"

func main() {
    // Init YANFF system
    config := flow.Config{}
    checkFatal(flow.SystemInit(&config))

    initCommonState()

    checkFatal(flow.SystemStart())
}

yanff-1$ cd $YANFF/examples/tutorial
yanff-1$ sudo ./step1

yanff-0$ cd $YANFF/examples/tutorial
yanff-0$ ./genscripts
yanff-0$ ./rumpktgen.sh

Pktgen:/> start 0
.....
Pktgen:/> quit
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

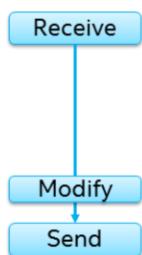
Optimization  
Notice



15

## Let's try (2 of 11)

Flow graph:



```
package main  
import "github.com/intel-go/yanff/flow"  
  
func main() {  
    config := flow.Config{}  
    checkFatal(flow.SystemInit(&config))  
  
    initCommonState()  
  
    firstFlow, err := flow.SetReceiver(0)  
    checkFatal(err)  
    checkFatal(flow.SetHandler(firstFlow, modifyPacket[0], nil))  
    checkFatal(flow.SetSender(firstFlow, 0))  
  
    checkFatal(flow.SystemStart())  
}
```

```
yanff-1$ sudo ./step2
```

```
yanff-0$ ./rumpktgen.sh  
Pktgen:/> load step2.pg  
Pktgen:/> start 0  
...  
Pktgen:/> quit
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

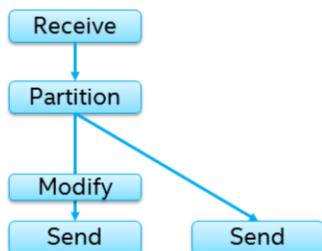


16

Example program receives packets from port 0 and sends them back to port 0. After “start 0” command pktgen should show that packets are not only sent on interface 0, but also the same number of packets is received back.

## Let's try (3 of 11)

Flow graph:



```
yanff-1$ sudo ./step3
```

```
yanff-0$ ./rumpktgen.sh
Pktgen:/> load step3.pg
Pktgen:/> start 0
```

```
...
Pktgen:/> quit
```

```
package main
import "github.com/intel-go/yanff/flow"

func main() {
    config := flow.Config{}
    checkFatal(flow.SystemInit(&config))

    initCommonState()

    firstFlow, err := flow.SetReceiver(0)
    checkFatal(err)
    secondFlow, err := flow.SetPartitioner(firstFlow, 300, 300)
    checkFatal(err)
    checkFatal(flow.SetHandler(firstFlow, modifyPacket[0], nil))
    checkFatal(flow.SetHandler(secondFlow, modifyPacket[1], nil))
    checkFatal(flow.SetSender(firstFlow, 0))
    checkFatal(flow.SetSender(secondFlow, 1))

    checkFatal(flow.SystemStart())
}
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

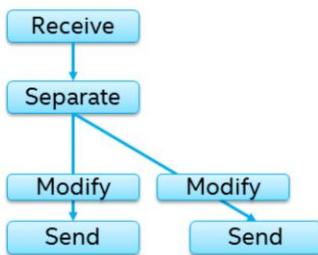


17

Example program receives packets from port 0 and sends them back half to port 0 and half to port 1. After “start 0” command pktgen should show that packets are sent on interface 0, and half of packets is received on interface 0, half on interface 1.

## Let's try (4 of 11)

Flow graph:



yanff-1\$ sudo ./step4

```
yanff-0$ ./rumpktgen.sh
Pktgen:/> load step4.pg
Pktgen:/> start 0
...
Pktgen:/> quit
```

```
package main
import "github.com/intel-go/yanff/flow"
import "github.com/intel-go/yanff/packet"

func main() {
    config := flow.Config{}
    checkFatal(flow.SystemInit(&config))
    initCommonState()

    firstFlow, err := flow.SetReceiver(0)
    checkFatal(err)
    secondFlow, err := flow.SetSeparator(firstFlow, mySeparator, nil)
    checkFatal(err)
    checkFatal(flow.SetHandler(firstFlow, modifyPacket[0], nil))
    checkFatal(flow.SetHandler(secondFlow, modifyPacket[1], nil))
    checkFatal(flow.SetSender(firstFlow, 0))
    checkFatal(flow.SetSender(secondFlow, 1))

    checkFatal(flow.SystemStart())
}

func mySeparator(cur *packet.Packet, ctx flow.UserContext) bool {
    cur.ParseL3()
    if cur.GetIPv4() != nil {
        cur.ParseL4ForIPv4()
        if cur.GetTCPForIPv4() != nil &&
            packet.SwapBytesUint16(cur.GetTCPForIPv4().DstPort) == 53 {
            return false
        }
    }
    return true
}
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

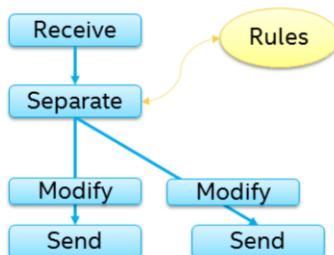


18

Example program receives packets from port 0 and sends back packets with port 53 to interface 1, all other packets to interface 0. Since step4.pg script generates packets with 11 different port numbers, 10/11 of packets should be received on port 0, and 1/11 should be received on port 1.

## Let's try (5 of 11)

Flow graph:



```
yanff-1$ sudo ./step5
```

```
yanff-0$ ./runpktgen.sh
Pktgen:/> load step5.pg
Pktgen:/> start 0
...
Pktgen:/> quit
```

```
import "github.com/intel-go/yanff/rules"
var L3Rules *rules.L3Rules

func main() {
    var err error
    config := flow.Config{}
    checkFatal(flow.SystemInit(&config))
    initCommonState()

    l3Rules, err = packet.GetL3ACLFromORIG("rules1.conf")
    checkFatal(err)

    firstFlow, err := flow.SetReceiver(0)
    checkFatal(err)
    secondFlow, err := flow.SetSeparator(firstFlow, mySeparator, nil)
    checkFatal(err)
    checkFatal(flow.SetHandler(firstFlow, modifyPacket[0], nil))
    checkFatal(flow.SetHandler(secondFlow, modifyPacket[1], nil))
    checkFatal(flow.SetSender(firstFlow, 0))
    checkFatal(flow.SetSender(secondFlow, 1))
    checkFatal(flow.SystemStart())
}

func MySeparator(cur *packet.Packet, ctx flow.UserContext) bool {
    return cur.L3ACLPermit(l3Rules)
}
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

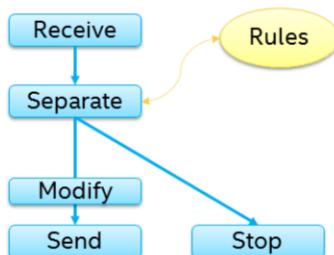


19

Example program receives packets from port 0 and sends back packets based on rules written in rules1.conf file. This rules file contents is written so that  $\frac{3}{4}$  of packets should be received on port 0 and  $\frac{1}{4}$  of packets should be received on port 1.

## Let's try (6 of 11)

Flow graph:



yanff-1\$ sudo ./step6

```
yanff-0$ ./runpktgen.sh
Pktgen:/> load step6.pg
Pktgen:/> start 0
...
Pktgen:/> quit
```

```
func main() {
    var err error
    config := flow.Config{}
    checkFatal(flow.SystemInit(&config))
    L3Rules = rules.GetL3RulesFromORIG("rules1.conf")
    checkFatal(err)
    firstFlow, err := flow.SetReceiver(0)
    checkFatal(err)
    secondFlow, err := flow.SetSeparator(firstFlow, mySeparator, nil)
    checkFatal(err)
    checkFatal(flow.SetHandler(firstFlow, modifyPacket[0], nil))
    checkFatal(flow.SetSender(firstFlow, 0))
    checkFatal(flow.SetStopper(secondFlow))
    checkFatal(flow.SystemStart())
}

func MySeparator(cur *packet.Packet, ctx flow.UserContext) bool {
    return cur.L3ACLPermit(l3Rules)
}
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

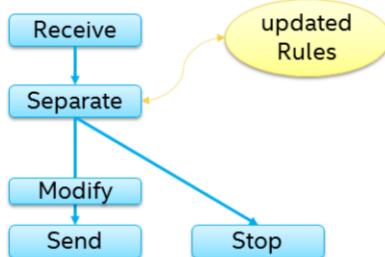


20

Example program receives packets from port 0 and sends back packets based on rules written in rules1.conf file. Only accepted packets are sent back to interface 0, rejected packets are dropped in this example. Rules file contents is written so that ¾ of packets should be received on port 0 and nothing should be received on interface 1.

## Let's try (7 of 11)

Flow graph:



yanff-1\$ sudo ./step7

yanff-0\$ ./runpktgen.sh  
Pktgen:/> load step7.pg  
Pktgen:/> start 0  
...  
Pktgen:/> quit

```
import "time"
var rulesp unsafe.Pointer
...
    l3Rules, err := packet.GetL3ACLFrom0RIG("rules1.conf")
    checkFatal(err)
    rulesp = unsafe.Pointer(&l3Rules)
    go updateSeparateRules()
...
func MySeparator(cur *packet.Packet, ctx flow.UserContext) bool {
    localL3Rules := (*packet.L3Rules)(atomic.LoadPointer(&rulesp))
    return cur.L3ACLPermit(localL3Rules)
}

func updateSeparateRules() {
    for {
        time.Sleep(time.Second * 5)
        localL3Rules, err := packet.GetL3ACLFrom0RIG("rules1.conf")
        checkFatal(err)
        atomic.StorePointer(&rulesp, unsafe.Pointer(localL3Rules))
    }
}
```

To make changes in rules1.conf file it is necessary to connect to target VM in another window or run YANFF executable in screen terminal multiplexer.

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

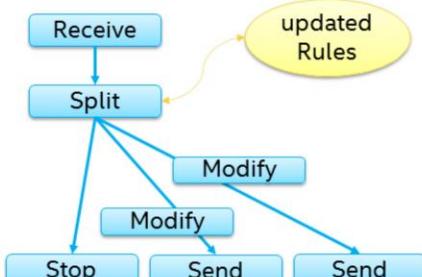


21

Example program receives packets from port 0 and sends back packets based on rules written in rules1.conf file. Only accepted packets are sent back to interface 0, rejected packets are dropped in this example. Rules file original contents is written so that  $\frac{3}{4}$  of packets should be received on port 0 and nothing should be received on interface 1. But this example allows on the fly modification of rules file, so for example changing Reject rule to "111.2.0.2/31" should change number of received packets on interface 0 to  $\frac{1}{2}$ .

## Let's try (8 of 11)

Flow graph:



yanff-1\$ sudo ./step8

```
yanff-0$ ./runkptgen.sh
Pktgen:/> load step8.pg
Pktgen:/> start 0
...
Pktgen:/> quit
```

```
const flowN = 3
firstFlow, err := flow.SetReceiver(0)
checkFatal(err)
outputFlows, err := flow.SetSplitter(firstFlow, mySplitter, flowN, nil)
checkFatal(err)
checkFatal(flow.SetStopper(outputFlows[0]))
for i := uint8(1); i < flowN; i++ {
    checkFatal(flow.SetHandler(outputFlows[i], modifyPacket[i-1], nil))
    checkFatal(flow.SetSender(outputFlows[i], i-1))
}

func mySplitter(cur *packet.Packet, ctx flow.UserContext) uint {
    localL3Rules := L3Rules
    return cur.L3ACLPort(localL3Rules)
}
```

To make changes in rules2.conf file it is necessary to connect to target VM in another window or run YANFF executable in screen terminal multiplexer.

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

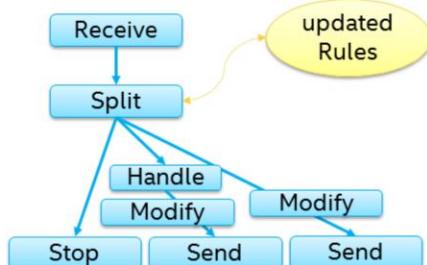


22

Example program receives packets from port 0 and sends back packets based on rules written in rules2.conf file. This rules file specifies to which output port the program should send a packet, port 0 meaning to drop packet. Rules file original contents is written so that  $\frac{1}{2}$  of packets should be received on port 0 and  $\frac{1}{4}$  should be received on port 1. But this example allows on the fly modification of rules file, so for example changing port 1 rule to "111.2.0.0/32" should change number of received packets on interface 0 to  $\frac{1}{4}$  and received packets on interface 1 to  $\frac{1}{2}$ .

## Let's try (9 of 11)

Flow graph:



yanff-1\$ sudo ./step9

```
yanff-0$ ./runkptgen.sh
Pktgen:/> load step9.pg
Pktgen:/> start 0
...
Pktgen:/> quit
```

```
import "github.com/intel-go/yanff/common"
firstFlow, err := flow.SetReceiver(0)
checkFatal(err)
outputFlows, err := flow.SetSplitter(firstFlow, mySplitter, flowN, nil)
checkFatal(err)
checkFatal(flow.SetStopper(outputFlows[0]))
checkFatal(flow.SetHandler(outputFlows[1], myHandler, nil))
for i := uint8(1); i < flowN; i++ {
    checkFatal(flow.SetHandler(outputFlows[i], modifyPacket[i-1], nil))
    checkFatal(flow.SetSender(outputFlows[i], i-1))
}

func myHandler(cur *packet.Packet, ctx flow.UserContext) {
    cur.EncapsulateHead(common.EtherLen, common.IPv4MinLen)
    cur.ParseL3()
    cur.GetIPv4NoCheck().SrcAddr = packet.BytesToIPv4(111, 22, 3, 0)
    cur.GetIPv4NoCheck().DstAddr = packet.BytesToIPv4(3, 22, 111, 0)
    cur.GetIPv4NoCheck().VersionIhl = 0x45
    cur.GetIPv4NoCheck().NextProtoID = 0x04
}
```

To make changes in rules2.conf file it is necessary to connect to target VM in another window or run YANFF executable in screen terminal multiplexer.

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

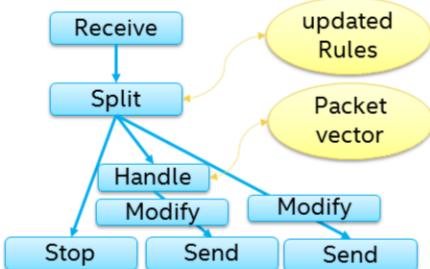


23

Example program receives packets from port 0 and sends back packets based on rules written in rules2.conf file. Difference with previous example is that port 1 flow also encapsulates a packet into IPv4 packet. Therefore received traffic on port 0 should be slightly higher than  $\frac{1}{2}$  of packets. If for example port 0 (drop) rule is changed to “111.2.1.2/32” (no packet will match)  $\frac{1}{2}$  of packets should be received on port 0 and  $\frac{1}{2}$  of packets should be received on port 1, but since on port 0 they are encapsulated, received traffic should be slightly higher on port 0.

## Let's try (10 of 11)

Flow graph:



```
func myHandler(curV []*packet.Packet, num uint, ctx flow.UserContext) {
    for i := uint(0); i < num; i++ {
        cur := curV[i]
        cur.EncapsulateHead(common.EtherLen, common.IPV4MinLen)
        cur.ParseL3()
        cur.GetIPv4NoCheck().SrcAddr = packet.BytesToIPv4(111, 22, 3, 0)
        cur.GetIPv4NoCheck().DstAddr = packet.BytesToIPv4(3, 22, 111, 0)
        cur.GetIPv4NoCheck().VersionIhl = 0x45
        cur.GetIPv4NoCheck().NextProtoID = 0x04
    }
}
```

```
yanff-1$ sudo ./step10
```

```
yanff-0$ ./runpktgen.sh
Pktgen:/> load step10.pg
Pktgen:/> start 0
...
Pktgen:/> quit
```

To make changes in rules2.conf file it is necessary to connect to target VM in another window or run YANFF executable in screen terminal multiplexer.

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice

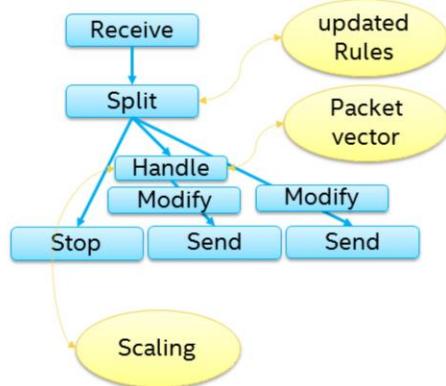


24

Example program receives packets from port 0 and sends back packets based on rules written in rules2.conf file. Difference with previous example is that handler function accepts a packet vector, everything else is the same.

## Let's try (11 of 11)

Flow graph:



To make changes in rules1.conf file it is necessary to connect to target VM in another window or run YANFF executable in screen terminal multiplexer.

```
func myHandler(curV []*packet.Packet, num uint, ctx flow.UserContext) {
    for i := uint(0); i < num; i++ {
        cur := curV[i]
        cur.EncapsulateHead(common.EtherLen, common.IPv4MinLen)
        cur.ParseL3()
        cur.GetIPv4NoCheck().SrcAddr = packet.BytesToIPv4(111, 22, 3, 0)
        cur.GetIPv4NoCheck().DstAddr = packet.BytesToIPv4(3, 22, 111, 0)
        cur.GetIPv4NoCheck().VersionIhl = 0x45
        cur.GetIPv4NoCheck().NextProtoID = 0x04
    }
    // Some heavy computational code
    heavyCode()
}
```

```
yanff-1$ sudo ./step11
```

```
yanff-0$ ./runpktgen.sh
Pktgen:/> load step5.pg
Pktgen:/> start 0
...
Pktgen:/> quit
```

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice



25

Example program receives packets from port 0 and sends back packets based on rules written in rules2.conf file. Difference with previous example is that vector handler function executes some heavy code to demonstrate automatic scaling. On systems with many CPU cores the example program should try to use more than in step 10 example.

## Finally: NAT

```
yanff-1$ ./genscripts -pktgen direct  
yanff-1$ sudo ../nat/main/nat -config nat.json
```

```
yanff-0$ ./runpktgen.sh  
Pktgen:/> load nat.pg  
Pktgen:/> start 0  
Pktgen:/> start 1  
...  
Pktgen:/> quit
```

Example demonstrates using NAT example with pktgen. Sample configuration for NAT configures port 0 as private network port and port 1 as public network port. Addresses are translated from 192.168.1.1/24 subnet to 10.1.1.1/24 subnet with NAT IP on public interface 10.1.1.1. Sample pktgen script implements 11 client addresses that try to access one server on public network and server sending packets back to the clients. When both “start 0” and “start 1” commands are executed pktgen should show the same traffic received on interface 0 and interface 1 which means that address translation works in both directions.

## Q & A ?

Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice 



27

# Optimization Notice

## Optimization Notice

Intel's compilers may or may not optimize to the same degree for non-Intel microprocessors for optimizations that are not unique to Intel microprocessors. These optimizations include SSE2®, SSE3, and SSSE3 instruction sets and other optimizations. Intel does not guarantee the availability, functionality, or effectiveness of any optimization on microprocessors not manufactured by Intel. Microprocessor-dependent optimizations in this product are intended for use with Intel microprocessors. Certain optimizations not specific to Intel microarchitecture are reserved for Intel microprocessors. Please refer to the applicable product User and Reference Guides for more information regarding the specific instruction sets covered by this notice.

Notice revision #20110804

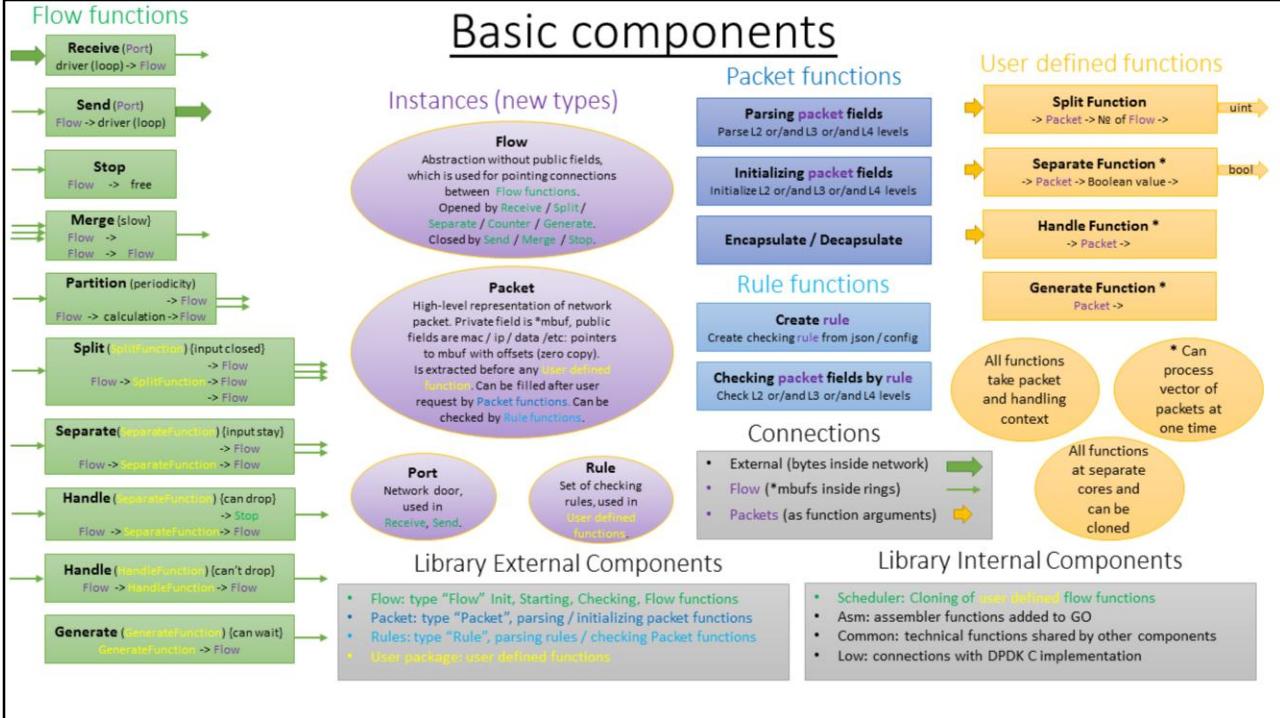
Copyright © 2017, Intel Corporation. All rights reserved.  
\*Other names and brands may be claimed as the property of others.

Optimization  
Notice



28

# Basic components



# Lab configuration



Jump host 207.108.8.161, Login: gashiman, Password: YanffLab

## Finally (2 of 2): ipsec

- Showing ipsec example